

24H

FIRST COMMIT → FINAL PUSH

Build a **slice** of the studio.

A 24-hour take-home for product-minded builders. We give you a fuzzy product and seven surface names — you decide what exists, what's faked, and what's worth doing beautifully.

ROLE · PRODUCT / FRONTEND / FULLSTACK

WINDOW · 24 HOURS

LANGUAGE · ENGLISH OR CHINESE

AI TOOLS · ENCOURAGED

01 The product

Heedo is building an **AI creative studio for e-commerce advertisers**. Teams bring product photos; the studio turns them into ad images and videos, keeps the results organized, and — eventually — learns from how those ads perform once they're live on real ad platforms.

The product is young. Parts of it exist, parts are half-decided, and parts are open questions we argue about at lunch. This brief is **deliberately vague**: turning a fuzzy product idea into something runnable is the job, and this assignment mirrors it.

02 The brief

Design and build a **working slice of this studio** — a web app we can clone, run, and click through end to end. Below are the surfaces and their intents. Scope, depth, data model, and visual language are yours to decide.

01

Gallery homepage

The front door. What greets a returning team: their recent work? curated picks? a reason to come back daily?

DEPTH · YOUR CALL

02

Templates

Nobody should start from a blank graph. A template is the DAG shape plus pre-filled values — e.g. a video template: two chained prompts (product pack shot → scenario image, scenario image → video) plus many params, ready to tweak.

DEPTH · YOUR CALL

03

Canvas

The heart of the app: a node canvas. Users compose a graph from many kinds of nodes, wiring them where types match — an image-generation node's output can feed an image-to-video node's input. AI generation may be mocked or real.

THE HEART

04

Credits

Generation costs money. Show the price before, charge after, and keep the balance legible at all times.

DEPTH · YOUR CALL

05

Team management

One workspace, several people. Members, roles, and the difference between who can spend and who can only look.

DEPTH · YOUR CALL

A coherent product where two surfaces go deep and three are honest stubs beats five half-screens.

CHOOSING WHAT TO CUT — AND TELLING US WHY — IS PART OF THE TEST

06

Ad platform connections

Meta / TikTok / YouTube. What does “connecting” an ad account look and feel like? What states can a connection be in? What syncs?

07

Open design question — performance data x canvas

Once ads run, numbers come back: impressions, CTR, spend, ROAS. How should that data show up *inside* the creative workspace, on or around the creatives themselves?

06 and 07 are open problems on our roadmap — there's no internal “right answer” we're comparing you against. Topic 07 may be **built, sketched, or argued in writing**: annotated screenshots, Figma frames, or a README essay all count. Strong reasoning here is worth more than pixels.

03 Ground rules

THE CLOCK

Your window is **24 hours** from your first commit to your final push — we read timestamps. We do not expect 24 hours of typing; sleep is allowed and recommended. We expect prioritization.

AMBIGUITY

If something feels underspecified, it is — intentionally. Make a call, write the assumption down, move on. Email only if truly blocked.

MOCK HONESTLY

Auth, payments, AI generation, ad APIs — all fakeable. Label the fakes and keep the seams clean, so we can see exactly where real systems would slot in.

AI TOOLS

Encouraged — Claude Code, Cursor, whatever you drive best. What we evaluate is the thinking that steered the tools, so show your steering (see deliverables).

STACK

Your choice. Frontend-only with mocked data is fine; a thin real backend is fine. For reference we run Next.js/TypeScript + FastAPI/Python — matching it is pleasant but earns no extra points.

04 Deliverables



A private GitHub repo

Invite `mianlang@devease.ai` and `x@devease.ai` as collaborators.



An honest git history

Commit as you go, with real messages — the log is the story of your 24 hours, and we read it. Don't squash, don't rewrite.



A README that does the talking

- How to run it — ideally ≤ 2 commands.
- What you built vs. cut, and why.
- Assumptions you made where the brief was vague.
- How you used AI tools — what worked, what you overrode, anything you're proud of or suspicious of.
- Your take on open question 07, if you didn't build it.
- "With one more week, I would ..."



Bonus, not required

A deployed link, or a screen recording (≤ 3 minutes).

05 How we read it

DIMENSION

WEIGHT

WHAT WE ASK OURSELVES

Delivery & coherence

 35%

Does it run? Does it feel like one product rather than disconnected screens?

Product judgment

 20%

Were the cut/deepen choices the right ones for a user? Do the flows make sense?

Process & reasoning

 20%

README, commit narrative, assumptions, how the AI tools were steered.

Craft

 15%

Visual and interaction quality where it matters most.

DIMENSION	WEIGHT	WHAT WE ASK OURSELVES
Frontier thinking	10%	The ad-platform connection and the performance-data-on-canvas exploration.

06 Fine print

- §

This is an evaluation exercise — we will **not** ship your code or designs.
- §

Write in **English or Chinese**, whichever lets you think faster.
- §

What happens next: we read everything (yes, the git log too) and get back to you over email. There's no live walkthrough round — your README and git history do the talking, so invest in them.
- §

Questions:

mianlang@devease.ai

.

Good luck — build something you'd be excited to demo. **We're rooting for you.**